

Artificial intelligence in education: applications, benefits and challenges – a systematic review

Kritica Jain
Department of Computer Science
Chitkara University
Chandigarh, India
kritica523.be22@chitkara.edu.in

Abstract—AI is everywhere now, and classrooms are no exception. There is a massive push towards smart tutoring, automated grading, and predictive analytics, fundamentally changing how educators teach and students learn. But despite all the hype, there is still a lot of confusion about what actually works in a real school setting versus what just looks good on paper. This review examines 15 peer-reviewed studies published between 2010 and 2022, drawn from IEEE Xplore, SpringerLink, and Google Scholar. Four key application areas are explored in depth: personalized learning through adaptive platforms, intelligent tutoring systems that simulate one-on-one instruction, automated assessment and grading tools, and learning analytics systems used for early student intervention. Key findings show that nearly half of existing AI in education research focuses on student profiling and performance prediction, while ethical frameworks remain almost entirely absent from the literature — only 5 out of 146 reviewed papers addressed ethics in any meaningful way. Research on AI in education also grew by over 400% between 2010 and 2020, yet mainstream classroom adoption remains limited due to infrastructure gaps, lack of teacher training, and unresolved concerns around data privacy, algorithmic bias, and transparency. Ultimately, institutions have to take responsibility for deploying these tools carefully and equitably, rather than adopting them blindly.

Keywords—Artificial intelligence, education technology, machine learning, adaptive learning, intelligent tutoring systems, learning analytics

I. INTRODUCTION

It is easy to see why schools are rushing to integrate AI. The traditional classroom model has always struggled with the reality of one teacher trying to manage thirty different learning paces. AI offers a scalable workaround. We already see this in action with platforms like Khan Academy or Duolingo, which use adaptive algorithms to change the difficulty of questions based on a user's previous answers. These are not futuristic concepts; they are already mainstream and used by millions of students globally every single day.

What makes this moment different from earlier waves of education technology is scale and sophistication. Earlier tools — overhead projectors, computer labs, online courses — changed the medium of instruction but not its logic. AI is different because it can respond to the individual. A student using an AI-powered platform gets a genuinely different experience from the student sitting next to them, based on their own history of responses, mistakes, and progress. That has never been possible before at this scale.

However, bringing this technology into formal education creates immediate complications. When a platform tracks every single keystroke a student makes, privacy becomes a serious liability. When algorithms start predicting which

students might fail, we have to look closely at the biases baked into those systems. A model trained on historical data from one type of institution will not necessarily make fair predictions for students from a completely different background.

There is also a glaring digital divide. When schools in rural areas or developing nations still struggle with basic internet access, AI-driven education risks becoming a luxury that widens the inequality gap rather than closing it. In India alone, a significant portion of students in rural government schools have no reliable access to the devices needed to benefit from any of these tools. The promise of AI in education cannot be separated from the infrastructure reality on the ground.

This paper takes stock of where things stand. Drawing on a systematic review of 15 published studies, it examines what AI is actually being used for in education, what the evidence says about its effectiveness, and what challenges still need to be worked through before AI can genuinely serve all learners — not just those in well-resourced institutions.

II. LITERATURE REVIEW

The academic literature on AI in education has grown rapidly over the past decade, spanning technical evaluations of specific tools, broad systematic reviews, and critical sociological analyses. This section reviews the most relevant and credible work across all three categories.

Holmes, Bialik, and Fadel provided one of the earlier comprehensive overviews of AI in education, focusing heavily on intelligent tutoring systems [1]. Their central argument was that AI has genuine potential to close learning gaps by giving every student the kind of one-on-one feedback that only a private tutor would normally provide. The evidence they presented was encouraging, though they were careful to note that most of the data came from controlled lab settings rather than real classrooms. A system that works perfectly in a research environment often performs very differently when deployed across a diverse student population.

When Zawacki-Richter and colleagues systematically reviewed 146 papers on AI in higher education, they found that 47.6% focused on student profiling and performance prediction [2]. Even more concerning — only 5 out of those 146 papers addressed ethical concerns in any meaningful way. Most research also came from just three countries: the USA, China, and the UK, raising serious questions about how applicable these findings are to other educational contexts including India.

Chen, Chen, and Lin reviewed AI applications across both K-12 and higher education and found that machine learning was the dominant technique used in the majority of studies examined [3]. They identified a significant lack of

longitudinal studies — meaning most research captures only short-term effects rather than long-term learning outcomes. A student might perform better on a test immediately after using an AI tutoring system, but whether that advantage holds six months later remains largely unanswered in the literature.

Luckin and colleagues made a point that is easy to overlook: the technology is only as useful as the people using it [4]. Their research found that institutional readiness and teacher training are often the real bottlenecks. In many deployments they studied, teachers were given tools without adequate training, leading to underuse, misuse, or outright abandonment of systems that might otherwise have been valuable.

Selwyn offers perhaps the most critical perspective in this literature [5]. His concern is that AI in education can easily become a tool of institutional surveillance rather than student support — tracking behavior, reducing teaching to data collection, and ultimately serving administrative interests more than student learning. He argues that the enthusiasm around AI in education often obscures a fundamental question: who actually benefits from all this data collection?

Zhai and colleagues reviewed AI in education research published between 2010 and 2020, finding that the number of studies grew by over 400% in that decade [6]. Machine learning and natural language processing accounted for the majority of techniques used, while deep learning was the fastest growing category. However, most studies were conducted in lab or controlled settings, and there was a notable lack of replication studies — meaning many findings have not been independently verified.

Woolf's work on intelligent tutoring systems provides important technical grounding for understanding why these systems work when they do [7]. His documentation of adaptive feedback mechanisms — the way ITS systems adjust hint specificity, problem difficulty, and feedback timing based on student responses — explains why students working with well-designed ITS often outperform those in standard instruction. Table I summarizes the key studies reviewed.

TABLE I. SUMMARY OF KEY LITERATURE ON AI IN EDUCATION

Author(s)	Application	Key finding
Holmes et al. [1]	Intelligent tutoring	Lab results good; real-world gaps remain
Zawacki-Richter et al. [2]	Adaptive learning	Only 5/146 papers on ethics
Chen et al. [3]	Predictive analytics	No longitudinal studies found
Luckin et al. [4]	Data-driven instruction	Teacher training is bottleneck
Selwyn [5]	Societal impact	AI risks becoming surveillance tool
Zhai et al. [6]	AI review 2010-2020	400% growth; mostly lab-based
Woolf [7]	Intelligent tutoring	Adaptive feedback drives gains

TABLE II. DISTRIBUTION OF AI RESEARCH FOCUS AREAS IN HIGHER EDUCATION [2]

Research focus area	Share of studies (%)
Student profiling and performance prediction	47.6%
Intelligent tutoring and adaptive learning	22.1%

Automated assessment and grading	17.3%
Learning analytics and dashboards	7.9%
Ethical and societal implications	3.4%
Other applications	1.7%

III. RESEARCH METHODOLOGY

To build this overview, a systematic literature review of peer-reviewed journals published over the last ten to fifteen years was conducted. The methodology follows a structured approach to ensure transparency and reproducibility in the selection and analysis of relevant studies.

The search was conducted across four major academic databases: IEEE Xplore, SpringerLink, Google Scholar, and ACM Digital Library. Search terms used included combinations of "Artificial Intelligence in Education," "intelligent tutoring systems," "adaptive learning platforms," "learning analytics," and "AI ethics education." The initial search returned over 200 potentially relevant results across all databases.

Inclusion criteria were applied to filter this initial pool. Papers were included if they: (1) focused on AI applications in real educational settings, (2) were published in peer-reviewed journals or refereed conference proceedings, (3) were published between 2010 and 2022, and (4) were available in English. Exclusion criteria removed papers that were purely theoretical, focused solely on non-educational AI, or lacked empirical or review-based findings.

This two-stage filtering process reduced the pool from over 200 papers to 15 highly relevant studies. These were then read in full and analyzed thematically around four areas: what AI is being used for, what benefits have been observed, what problems keep appearing, and what ethical issues researchers are raising.

This approach has limitations worth acknowledging. It depends on what got published and what got indexed in English-language databases. Research published in other languages or in regional journals not covered by these databases is not represented, meaning findings skew toward research from English-speaking and technologically advanced countries. Despite this, the review provides a solid overview of the mainstream academic literature on AI in education.

IV. APPLICATIONS OF ARTIFICIAL INTELLIGENCE IN EDUCATION

AI is entering classrooms in several specific ways. Four of the most significant and well-documented applications are discussed in detail below.

A. Personalized learning

The core concept behind personalized learning is not new — teachers have always known that students learn at different speeds. What AI adds is the ability to act on that knowledge at scale, in real time. Platforms like Knewton and Carnegie Learning analyze how a student is performing moment by moment and modify what content they see next accordingly. If someone keeps getting fraction problems wrong, the system gives them more targeted practice before moving on. Research consistently finds that this kind of adaptive instruction keeps students more engaged and improves final scores [6]. The effect is strongest in subjects where skills build sequentially — mathematics, language learning, and coding.

What separates AI-driven personalization from earlier adaptive tools is the depth of the feedback loop. Earlier systems adjusted difficulty based on simple right/wrong signals. Modern AI platforms incorporate natural language processing, response time analysis, and pattern recognition across thousands of past learners to make far more nuanced adjustments. This means a student's learning path genuinely diverges from their peers in ways that reflect their individual strengths and gaps rather than just their overall score.

B. Intelligent tutoring systems

Intelligent Tutoring Systems (ITS) go further than adaptive content delivery. Rather than simply adjusting what material a student sees, they try to replicate the experience of working with a human tutor — asking guiding questions, offering hints when a student is stuck, and explaining what went wrong rather than just marking answers incorrect. Systems like AutoTutor have been tested in university settings across subjects including physics, biology, and reading comprehension, sometimes achieving learning gains comparable to human tutoring [7].

That is a genuinely striking finding, though it needs context. ITS tend to work best in well-defined domains where correct answers are clear and feedback can be precise. For subjects requiring creative thinking, ethical reasoning, or nuanced argumentation — areas central to humanities and social sciences — human tutors still have a significant advantage. The technology is powerful within its limits, but recognizing those limits is essential to deploying it responsibly.

C. Automated assessment and grading

Grading is one of the most time-consuming parts of teaching, and it is an area where AI is increasingly making a practical difference. Natural language processing tools can now evaluate short answers, coding assignments, and even essays with reasonable accuracy. Tools like Turnitin's AI grading features are already widely used across university campuses. Ifenthaler and Yau found that institutions using data-driven assessment tools reported not just faster feedback cycles but also measurably higher student satisfaction scores compared to traditional grading [8].

However, accuracy varies significantly by task type. For structured tasks with clear rubrics, AI grading performs consistently well. For assignments requiring nuanced judgment about argument quality or originality, human oversight remains essential. The technology works best as a first pass that reduces teacher workload, not as a final arbiter of student performance.

D. Learning analytics and early intervention

Predictive learning analytics may currently be the most practically valuable application of AI in higher education. These systems monitor patterns in student engagement — attendance, assignment timing, login frequencies, quiz scores, forum activity — to build a behavioral profile for each student. When the system detects a pattern associated with disengagement, it flags the student so an advisor can reach out proactively. D'Mello and Graesser found that students who received timely alerts and targeted support were significantly less likely to disengage completely from their courses [9].

Universities that have implemented these early-warning systems consistently report improvements in retention rates. Identifying a struggling student in week three rather than week

ten dramatically increases the chances of a successful intervention. The data required to run these systems is already being collected by most modern learning management systems — the main challenge is putting it to use effectively. Table II gives an overview of key benefits and challenges, and Table IV shows the growth of AI research in education over the past decade.

TABLE III. BENEFITS AND CHALLENGES OF AI IN EDUCATION

Benefits	Challenges
Personalized learning at scale	Student data privacy concerns
Faster, automated assessment	Algorithmic bias and unfair outcomes
Early identification of at-risk students	High infrastructure and implementation cost
Increased student engagement	Lack of transparency in AI decisions
Reduced teacher administrative burden	Reduced meaningful human interaction
Scalable quality education delivery	Dependence on internet and devices

TABLE IV. GROWTH OF AI IN EDUCATION RESEARCH 2010-2020 [6]

Period	Approx. no. of studies	Dominant technique
2010-2012	~45	Rule-based systems
2013-2015	~110	Machine learning
2016-2017	~180	Machine learning + NLP
2018-2019	~290	Deep learning emerging
2020	~380	Deep learning + NLP

V. CHALLENGES AND ETHICAL CONCERNS

The benefits described above are real, but so are the problems. Four challenges stand out as particularly serious and consistently underaddressed in the literature.

Data privacy and security: every interaction a student has with an AI-powered platform generates data. Over time this builds into a highly detailed profile covering academic strengths, weaknesses, work habits, and engagement patterns. The question of who owns this data, who can access it, and what happens to it after graduation is rarely answered clearly. The UNESCO policy report found that as of 2021, over 40 countries had developed national AI strategies, yet fewer than 10% included specific guidelines for AI use in education [10]. In India, data protection legislation is still catching up with ed-tech reality, leaving students to click "agree" on terms of service they do not fully understand.

Williamson and Eynon traced this problem back further, arguing that surveillance and data exploitation in education are not new — they are longstanding institutional practices that AI has made faster, more granular, and harder to contest [11]. Their historical analysis is a useful reminder that the ethical challenges around AI in education are not purely technical problems. They reflect deeper questions about power and accountability that better engineering alone will not resolve.

Algorithmic bias and fairness: algorithms are only as objective as the data they were trained on, and training data almost always reflects existing inequalities. A predictive model built primarily on data from well-resourced urban universities will likely produce unfair predictions for students

from rural or lower-income backgrounds. Bond and colleagues found that even in well-regulated German universities, student awareness of how their data is collected remained consistently low [12]. If students do not know their data is being used to make predictions about their futures, meaningful informed consent is impossible.

Transparency and the black-box problem: most deep learning systems cannot explain their own decisions in any meaningful way. If an AI system flags a student as at risk of failing, that student deserves to know exactly why — what data was used, what patterns were detected, and what assumptions the model is making. Without this transparency, students cannot challenge incorrect predictions, and teachers cannot make informed decisions about when to trust the system and when to override it.

Over-reliance and deskilling: Roll and Wylie found that despite decades of ITS research, mainstream classroom adoption remains surprisingly limited — largely because teachers feel excluded from the design process and uncertain about how to use these tools effectively [13]. The most advanced AI tool is useless if teachers do not trust or understand it. Beyond adoption, there is also the longer-term risk that over-reliance on automation erodes the professional judgment of educators themselves. Teaching is a relational practice built on human connection and contextual judgment — things that cannot and should not be automated.

VI. RESULTS AND DISCUSSION

The review of 15 studies reveals several consistent and important patterns worth drawing out explicitly.

First, AI applications in education are heavily concentrated in a small number of areas. Zawacki-Richter and colleagues found that nearly half of all research — 47.6% — focuses on student profiling and performance prediction alone [2], as shown in Table III. Personalized learning and intelligent tutoring account for much of the rest. This concentration means some potentially important applications — AI for teacher development, curriculum design, and accessibility — are significantly underresearched.

Second, the geographic distribution of this research is strikingly uneven. The vast majority of studies come from the USA, China, and the UK [2]. South Asia, Sub-Saharan Africa, and Latin America are barely represented. A tool validated at a well-funded American university cannot simply be assumed to work in a government school in rural India. Educational systems, student demographics, and infrastructure availability vary enormously across regions.

Third, ethical considerations remain an afterthought. Only 5 of 146 papers in Zawacki-Richter's sample addressed ethics meaningfully [2]. The research community is producing tools and evaluating technical performance far faster than it is asking whether those tools should be deployed and under what conditions. This is a serious governance gap that needs urgent attention.

Fourth, despite over 400% growth in research between 2010 and 2020 — as shown in Table IV — mainstream classroom adoption remains limited [6]. Roll and Wylie found that cost, infrastructure barriers, and failure to involve teachers in design are the primary reasons [13]. D'Mello and Graesser's findings on early intervention offer some optimism — institutions that deploy these systems well see real retention

improvements [9] — but these remain the exception rather than the rule.

Finally, almost all of this literature is based on short-term studies in controlled settings. Chen and colleagues identified the lack of longitudinal research as a major gap [3]. We do not yet have strong evidence about long-term effects of AI-mediated education on student learning, motivation, or skill development — a significant limitation on how confident we can be in any of the positive findings reported.

VII. CONCLUSION

AI is not going to solve all of education's problems, and anyone who claims otherwise is overselling it. But the evidence reviewed here does suggest that, used thoughtfully and with proper institutional support, these tools can make real differences — particularly for students who currently lack access to individualized attention.

The most consistent finding across this review is not about any particular technology but about context. Tools work when teachers are trained and supported. They work when infrastructure is reliable. They work when the people designing them have actually talked to the students and teachers who will use them. Chiu found that student engagement with AI tools depends heavily on autonomy and competence — not just content quality — meaning that how tools are introduced matters as much as the tools themselves [14].

Garcia-Penalvo similarly argued that for AI to genuinely transform education, institutions need to move beyond isolated pilot projects and commit to systemic integration backed by clear policy frameworks [15]. That recommendation applies with particular force in countries like India, where the gap between the possibilities of AI in education and the current infrastructure reality is enormous. The technology will not close that gap on its own. Policy, investment, and genuine commitment to equity are what will.

The most pressing recommendation from this review is straightforward: before deploying AI in any classroom, institutions should ask not just whether the technology works, but who it works for, who it might harm, and whether students had any say in the decision. That is not a technical question. It is a question of values — and it deserves a serious answer.

ACKNOWLEDGMENT

The author thanks the faculty mentors at the Department of Computer Science, Chitkara University, for their guidance and support throughout this research work.

REFERENCES

- [1] W. Holmes, M. Bialik, and C. Fadel, *Artificial intelligence in education: promises and implications for teaching and learning*. Boston: Center for Curriculum Redesign, 2019.
- [2] O. Zawacki-Richter, V. I. Marin, M. Bond, and F. Gouverneur, "Systematic review of research on artificial intelligence applications in higher education," *Int. J. Educ. Technol. High. Educ.*, vol. 16, no. 1, pp. 1-27, 2019.
- [3] L. Chen, P. Chen, and Z. Lin, "Artificial intelligence in education: a review," *IEEE Access*, vol. 8, pp. 75264-75278, 2020.
- [4] R. Luckin, W. Holmes, M. Griffiths, and L. B. Forcier, *Intelligence unleashed: an argument for AI in education*. London: Pearson Education, 2016.
- [5] N. Selwyn, *Should robots replace teachers? AI and the future of education*. Cambridge: Polity Press, 2019.

- [6] X. Zhai et al., "A review of artificial intelligence (AI) in education from 2010 to 2020," *Complexity*, vol. 2021, pp. 1-18, 2021.
- [7] B. P. Woolf, *Building intelligent interactive tutors: student-centered strategies for revolutionizing e-learning*. Burlington: Morgan Kaufmann, 2010.
- [8] D. Ifenthaler and J. Y. K. Yau, "Utilising learning analytics to support study success in higher education," *Educ. Technol. Soc.*, vol. 23, no. 1, pp. 1-13, 2020.
- [9] S. K. D'Mello and A. Graesser, "AutoTutor and affective AutoTutor: learning by talking with intelligent computers," *ACM Trans. Interact. Intell. Syst.*, vol. 2, no. 4, pp. 1-39, 2012.
- [10] UNESCO, *Artificial intelligence in education: guidance for policy-makers*. Paris: UNESCO, 2021.
- [11] B. Williamson and R. Eynon, "Historical threads, missing links, and future directions in AI in education," *Learn. Media Technol.*, vol. 45, no. 3, pp. 223-235, 2020.
- [12] M. Bond, V. I. Marin, C. Dolch, S. Bedenlier, and O. Zawacki-Richter, "Digital transformation in German higher education," *Int. J. Educ. Technol. High. Educ.*, vol. 15, no. 1, pp. 1-20, 2018.
- [13] J. Roll and R. Wylie, "Evolution and revolution in artificial intelligence in education," *Int. J. Artif. Intell. Educ.*, vol. 26, no. 2, pp. 582-599, 2016.
- [14] T. K. F. Chiu, "Applying the self-determination theory (SDT) to explain student engagement in online learning during COVID-19," *J. Res. Technol. Educ.*, vol. 54, no. 1, pp. 14-30, 2022.
- [15] F. J. Garcia-Penalvo, "Artificial intelligence in education: exploring the potential for educational innovation," *Sustainability*, vol. 13, no. 18, Art. no. 10038, 2021.

